

## Automated EBS Backup & Recovery System

### Project Overview

**Services used:** EC2, EBS, Lambda, CloudWatch Events (EventBridge), IAM, CloudWatch Logs, (optional: SNS)

**Goal:** Automate EBS snapshot creation on a schedule, tag snapshots for identification, and implement a retention policy that deletes snapshots older than a configured retention window. Ensure reliable disaster recovery for compute workloads with minimal operational effort.

**Skills demonstrated:** Cloud automation, backup strategies, AWS Lambda scripting (Python / boto3), event-driven architecture, IAM least-privilege, monitoring.

### Architecture (high level)

1. **EventBridge (CloudWatch Events)** triggers an AWS Lambda function on a schedule (e.g., daily at 02:00 UTC).
  2. The **Lambda function** discovers EBS volumes to back up (via tags or by scanning instances), creates snapshots, and tags the snapshots with metadata (Name, BackupType, CreatedBy, RetentionDays, OriginVolume, InstanceID).
  3. The Lambda function also **applies retention rules**: it lists snapshots that match the tag and deletes those older than the configured retention period.
  4. **CloudWatch Logs** capture Lambda logs. Optionally, **SNS** notifies administrators of success/failure.
- 

### Design decisions & assumptions

- Volumes to be backed up are identified by a tag (recommended: Backup=true or Backup: daily) — this keeps scope explicit and safe.
  - Snapshots are tagged with the creation timestamp and retention window to make deletion deterministic.
  - Retention is enforced by a single Lambda (create + cleanup). Alternatively, separate Lambdas could be used for separation of concerns.
  - For consistency of application-level data, you may wish to coordinate with the instance OS (freeze writes, flush fs) — for Linux, use aws ec2 create-image or application-specific hooks. This project focuses on EBS-level snapshots.
- 

### Prerequisites

1. AWS account with permission to create Lambda, EventBridge Rule, IAM Role, and to create/delete EBS snapshots.
  2. Python 3.11 runtime for Lambda (or Node.js if you prefer — sample uses Python).
  3. AWS CLI / Console access to deploy and test.
- 

### Security & IAM

**IAM Role for Lambda** should follow least privilege. Example permissions (scope down with Resource ARNs):

- ec2:DescribeVolumes
- ec2:DescribeInstances
- ec2>CreateSnapshot

- ec2:CreateTags
- ec2:DescribeSnapshots
- ec2>DeleteSnapshot
- logs>CreateLogGroup, logs>CreateLogStream, logs>PutLogEvents
- sns>Publish

```
{  
    "Version": "2012-10-17",  
  
    "Statement": [  
        {  
            "Sid": "EC2SnapshotAndDescribePermissions",  
            "Effect": "Allow",  
            "Action": [  
                "ec2:DescribeVolumes",  
                "ec2:DescribeInstances",  
                "ec2>CreateSnapshot",  
                "ec2>CreateTags",  
                "ec2:DescribeSnapshots",  
                "ec2>DeleteSnapshot"  
            ],  
            "Resource": "*"  
        },  
        {  
            "Sid": "CloudWatchLogsWrite",  
            "Effect": "Allow",  
            "Action": [  
                "logs>CreateLogGroup",  
                "logs>CreateLogStream",  
                "logs>PutLogEvents"  
            ],  
            "Resource": "*"  
        }  
    ]  
}
```

```

        "Resource": "*"
    }
]
}

```

## STEP 1 — Tag the EBS Volumes for Backup

1. Open the AWS Console → Go to **EC2 Dashboard**.
2. In the left menu, click **Elastic Block Store** → **Volumes**.
3. Select the volume(s) you want to back up automatically.
4. Click **Tags** → **Add Tag**.
5. Enter:
  - **Key:** Backup
  - **Value:** daily
6. Click **Save**.

Volume details | EC2 | us-east-1

us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#VolumeDetails:volumId=vol-0f590033009c39904

aws EC2 > Volumes > vol-0f590033009c39904

Last updated 1 minute ago

**Actions** | **Delete** | **Modify**

Details	Value	Details	Value
Volume ID	vol-0f590033009c39904	Size	100 GiB
AWS Compute Optimizer finding	(i) Opt-in to AWS Compute Optimizer for recommendations.   Learn more	Type	gp3
Fast snapshot restored	No	IOPS	3000
Attached resources	-	Availability Zone	use1-az1 (us-east-1a)
Source	-	Created	Sat Nov 29 2025 12:12:10 GMT+0530 (India Standard Time)
Snapshot ID	-	Outposts ARN	-
		Managed	false
		Operator	-

aws EC2 > Volumes > vol-0f590033009c39904

United States (N. Virginia) Account ID: 1623-0849-0504 kavyanjali

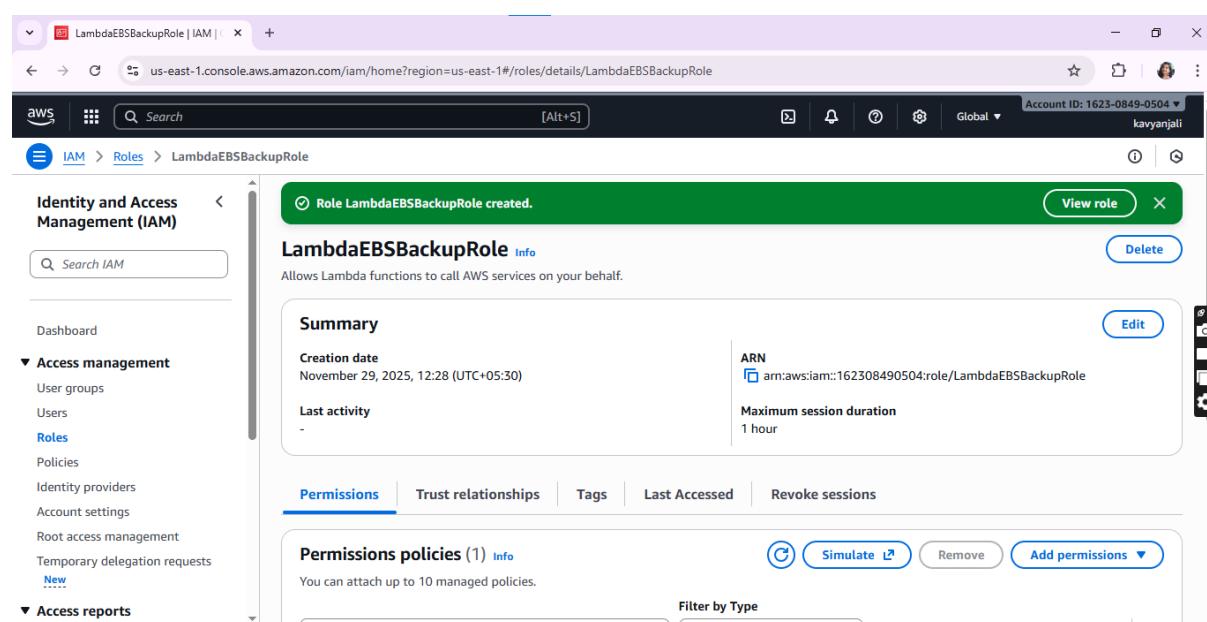
Status checks | Monitoring | **Tags**

Manage tags

Key	Value
Backup	daily

## STEP 2 — Create the IAM Role for the Lambda Function

1. Go to **IAM Console**.
2. Click **Roles** → **Create role**.
3. Select **AWS Service** → Choose **Lambda** → Click **Next**.
4. Click **Create Policy** (opens in new tab):
  - Choose **JSON** tab.
  - Paste the minimum required permissions (EC2 Describe/Create/Delete, Logs permissions).
  - Click **Next** → Name it: **EBSBackupPolicy** → Create.
5. Return to role creation window:
  - Search and select the **EBSBackupPolicy**.
  - Add **AWSLambdaBasicExecutionRole**.
6. Name the role: **LambdaEBSBackupRole**.
7. Click **Create Role**.



## STEP 3 — Create the Lambda Function

1. Open **AWS Lambda Console**.
2. Click **Create function**.
3. Choose:
  - **Author from scratch**
  - Function name: **EBS-Auto-Backup**
  - Runtime: **Python** (choose the latest available)
  - Execution role: **Use existing role** → select **LambdaEBSBackupRole**
4. Click **Create Function**.

### Upload Code

1. **Code Source** section.

```
import os
import boto3
import logging
from datetime import datetime, timezone, timedelta
```

```

logger = logging.getLogger()
logger.setLevel(logging.INFO)

ec2 = boto3.client('ec2')
sns = boto3.client('sns') if os.getenv('SNS_TOPIC_ARN') else None

BACKUP_TAG_KEY = os.getenv('BACKUP_TAG_KEY', 'Backup')
BACKUP_TAG_VALUE = os.getenv('BACKUP_TAG_VALUE', 'daily')
RETENTION_DAYS = int(os.getenv('RETENTION_DAYS', '7'))
CREATED_BY_TAG = 'ebs-automated-backup'
SNAPSHOT_DESC_PREFIX = os.getenv('SNAPSHOT_DESCRIPTION_PREFIX',
'AutoBackup')

def find_volumes_to_backup():
    # Find volumes tagged with BACKUP_TAG_KEY=BACKUP_TAG_VALUE
    filters = [
        {'Name': f'tag:{BACKUP_TAG_KEY}', 'Values': [BACKUP_TAG_VALUE]}
    ]
    resp = ec2.describe_volumes(Filters=filters)
    volumes = resp.get('Volumes', [])
    logger.info(f"Found {len(volumes)} volumes to back up")
    return volumes

def create_snapshot(volume):
    vol_id = volume['VolumeId']
    timestamp = datetime.now(timezone.utc).isoformat()
    desc = f'{SNAPSHOT_DESC_PREFIX} {vol_id} {timestamp}'
    resp = ec2.create_snapshot(VolumeId=vol_id, Description=desc)
    snap_id = resp['SnapshotId']
    tags = [
        {'Key': 'CreatedBy', 'Value': CREATED_BY_TAG},
        {'Key': 'RetentionDays', 'Value':

```

2. Replace the default code with the provided Lambda script.
3. Click **Deploy**.

#### Add Environment Variables

1. In the Lambda page → Click **Configuration**.
2. Choose **Environment variables** → **Edit** → **Add environment variable**.
  - BACKUP\_TAG\_KEY → Backup
  - BACKUP\_TAG\_VALUE → daily
  - RETENTION\_DAYS → 7
  - (Optional) SNS\_TOPIC\_ARN
3. Click **Save**.

## Set Timeout

1. Go to **Configuration** → **General configuration** → **Edit**.
2. Change Timeout to **1–3 minutes**.
3. Click **Save**.

The screenshot shows the AWS Lambda Functions console. The function name is "EBS-Auto-Backup". In the "Function overview" section, there is a "Diagram" button, a "Template" button, and a "Layers" section showing "(0) Layers". Below the diagram, there are buttons for "+ Add trigger" and "+ Add destination". On the right side, there is a "Description" field with a minus sign, a "Last modified" field showing "19 seconds ago", a "Function ARN" field with the value "arn:aws:lambda:us-east-1:162308490504:function:EBS-Auto-Backup", and a "Function URL" field with a "Info" link.

## STEP 4 — Create a Scheduled Event (EventBridge Rule)

1. Go to **Amazon EventBridge Console**.
2. Click **Rules** → **Create rule**.
3. Rule name: DailyEBSBackupSchedule.
4. Choose **Schedule**.
5. Enter CRON expression for daily 2 AM:

The screenshot shows the AWS EventBridge Rules console. The rule name is "DailyBackupTrigger". In the "Rule details" section, the rule name is "DailyBackupTrigger", the status is "Enabled", the event bus name is "default", the event bus ARN is "arn:aws:events:us-east-1:162308490504:rule/DailyBackupTrigger", and the type is "Standard". Below this, there are tabs for "Event pattern", "Targets", "Monitoring", and "Tags". The "Event pattern" tab shows the CRON expression: "1 { 2 "source": ["aws.scheduler"] 3 }".

## Test the Lambda function manually

1. Lambda console → select EBS-Auto-Backup → **Test**.
2. **Configure test event** → name it e.g., manual-test → leave default template (or empty JSON {}) → **Create**.

3. Click **Test**.
4. Immediately check **Execution result** (top-right) and **Logs** in the console.
5. Console → **Services** → **EC2** → **Snapshots** → confirm new snapshot(s) appear with tags:
  - CreatedBy: ebs-automated-backup
  - RetentionDays and BackupTimestamp

## Monitoring & Alerts

### CloudWatch Logs

1. Console → **CloudWatch** → **Log groups** → select /aws/lambda/EBS-Auto-Backup.
2. Inspect recent log streams for successful runs and errors.

The screenshot shows the AWS CloudWatch Log Events interface. The left sidebar lists various monitoring services: Favorites and recents, Alarms, AI Operations, GenAI Observability, Application Signals (APM), and Infrastructure Monitoring. The main area displays a log event from the 'aws-waf-logs-CloudFrontDistribution-EAFEBWWF99HM' log group. The event details are as follows:

Timestamp	Message
2025-11-24T04:03:30.086Z	No older events at this moment. <a href="#">Retry</a> {"timestamp":1763957010086,"formatVersion":1,"webaclId":"arn:aws:wafv2:us-east-1:162308490504:global/webacl/..."} No newer events at this moment. Auto retry paused. <a href="#">Resume</a>

## Check snapshots:

Go to EC2 → Snapshots.

Verify that new snapshots are created daily.

Check tags: CreatedBy, DeleteOn.

Verify old snapshots disappear after retention period.

## **Restore from Snapshot (Recovery)**

If your volume fails or you need to roll back:

1. Go to EC2 → Snapshots.
2. Choose your snapshot → Actions → Create Volume from Snapshot.
3. Choose the same Availability Zone as your instance.
4. Once created, Attach Volume to your EC2 instance.
5. Mount it inside the instance to access data.

## Final Summary – Automated EBS Backup & Retention System (AWS)

This project implements a fully automated backup and retention mechanism for Amazon EBS volumes using **Lambda**, **CloudWatch Events**, **EC2**, and **EBS Snapshots**. It removes the need for manual backups and ensures reliable disaster recovery for EC2 workloads.